

# Discerning Performance, Power, Energy and Area Efficacies of Democratized ISA Effort

## ABSTRACT

Instruction Set Architecture (ISA) is fundamental to how a wide variety of modern computer systems – ranging from simple handheld mobile devices to large scale data centers and server farms – are conceived, designed and implemented. ISA designers’ goal is often to capture the most basic functions and tasks that can be then used to compose and express complex applications and softwares. The General expectation is that a computing system should perform functions captured by ISA set in most possible efficient way in terms of performance (single-cycle) and power/energy.

While an ISA is central to computer design, there have been only a handful of successful ISAs forcing designers to choose them from a small subset even though it might not be the most efficient in capturing higher level applications. Unlike compilers, OSs, drivers, and other software components, ISAs have been a proprietary component by-and-large.

Democratization of ISA was the main theme behind the advent of RISC-V. It was touted to relieve the designer community and small to mid-scale OEMs from the clutches of proprietary ISA suppliers. While this is a novel thought in spirit, much depends upon the efficacies of democratized ISA itself. In this work, we set out to discern and quantify the viability of an open source ISA such as RISC-V.

## 1. INTRODUCTION

Numerous isas have been designed but couldn’t survive because they either didn’t offer anything unique or they were poorly designed.

RISCV is an emerging Instruction Set Architecture (ISA) and has become an important option for both academia and industry when considering new microprocessor designs. Features like modularity, extensibility, simplicity, and being open and free to use, make RISCV an attractive option for next generation of processors especially in embedded systems domain where new, customized, low-power, and efficient cores are needed. In this paper we present a comparative study on impacts of three well-known Instruction Set Architectures (ISAs) (MIPS, ARM, and RISCV) on Performance, Power, and Area (PPA) of state-of-the-art embedded processors through a systematic measurement campaign using several different toolchains and frameworks and sev-

eral standard benchmark suites. We particularly study the impact of these ISAs on important metrics such as static and dynamic Instruction Count (*icount*), Cycle Count, Microarchitectural Statistics (e.g. MPKI, Branch Prediction Accuracy, etc.), Dynamic Power, and Core’s Area and report our key findings on impacts of using different ISAs on each of these metrics. We find that some of these metrics are ISA-dependent and others are dependent on other factors such as compiler, runtime libraries, and specific microarchitectural features. Our main conclusion is that while comparing to MIPS and ARM, RISCV has some shortcomings and design/toolchain issues that should be addressed and fixed, due to its intrinsic features such as modularity it provides a great opportunity for designing customized PPA-efficient cores.

Instruction Set Architectures (ISAs) has a key role in designing cores for different domains, where x86 ISA has become dominant in desktop and server domains, and ARM has become the dominant ISA in mobile, tablet, and embedded system domain. The question of impact of ISA design on different Performance, Power, Area (PPA) metrics has traditionally been an important concern for designers and semiconductor industry especially in the 1980s and 1990s when chip area and processor design complexity were the primary constraints [24, 12, 17, 7]. In the past decade, we radical changes in computing landscape and rise of mobiles and tables and increasing popularity of ARM ISA this question again becomes an important issue.

Today, with proliferation of embedded and cyber-physical systems (e.g. IoTs) and increasing popularity of domain-specific languages and emerging applications like machine-learning and more importantly, introduction of a new, open-source, modular ISA (RISCV), this question once again becomes an interesting topic for research.

To answer this question, in this paper we present a comparative study on impacts of using three different ISAs (MIPS, ARM, and RISCV) on important metrics such as static and dynamic instruction count (*icount*), total execution time (cycle count), dynamic power, and area. We show which of these metrics are ISA-dependent and what are the other important factors on PPA. Using these experiments we pinpoint the shortcomings, issues, and advantages of using RISCV ISA over ARM

and MIPS ISAs.

## 2. BACKGROUND

RISC-V is an emerging open-source software and hardware ecosystem that has gained in popularity in both industry and academia [2, 11]. At the heart of the ecosystem, the RISC-V ISA is designed to be open, simple, extensible, and free to use. The RISC-V software tool chain includes open-source compilers (e.g., GNU/GCC and LLVM), a full Linux port, a GNU/GDB debugger, verification tools, and simulators. On the hardware side, several RISC-V prototypes (e.g., Celerity [4]) have been published. The rapid growth of the RISC-V ecosystem enables computer architects to quickly leverage RISC-V in their research.

## 3. METHODOLOGY

To study the effects of ISAs on Power, Performance, and Area, we used several different metrics using 4 different tools and more than 12 standard benchmark applications. Followings describe the frameworks, metrics, and benchmarks used in this paper. The reader can skip this section if he/she is uninterested in these details.

### 3.1 ISAs and Compiler

Table 1 shows the ISAs used in study. For each of these ISAs, we use *gnu-gcc* cross-compiler. We intentionally chose gcc so that we can use the same front-end to generate all binaries. All target independent optimizations are enabled (O3); machine specific tuning is disabled so there is a single set of binaries for each ISA. For MIPS, *-march* is used to generate MIPS release 6, 32bit and 64bit versions. For ARM, two separate compilers ARMv7 and AARCHv8 are used to generate 32bit and 64bit ARM binaries. Finally, for RISC-V we use the gnu-toolchain provided by RISC-V developers publicly available in github. We believe using similar flags and front-end could help us to mitigate the effect of compilers on performance and power metrics, however, we will later show that RISC-V compiler does have important inefficiencies and issues that could hurt the performance of the system.

Table 1: ISAs used in this study.

ISAs	Specification
ARM	32v7, 64v8 (AARCH)
MIPS	32r6, 64r6
RISC-V	rv32g (IMAFD), rv64g

### 3.2 Framework

#### 3.2.1 QEMU

To find the dynamic and static instruction count (ICOUNT), we use a well-known open-source emulator called Quick-EMULATOR (QEMU). QEMU is a hosted virtual machine monitor: it emulates the machine’s processor through

dynamic binary translation and provides a set of different hardware and device models for the machine, enabling it to run a variety of guest operating systems. We chose QEMU primarily cause it can emulate MIPS, ARM, and RISC-V ISAs in user-mode.

**Static Instruction Count:** Static icount is a classic metric to show the code density of different ISAs which can directly affect the performance, power, and area (i.e. required icache size). While static icount can simply be measured by measuring the lines of assembly code in a binary, a more meaningful and useful way to measure this metric is to count number of unique PCs (each PC represents an instruction) seen during the execution of the an application (i.e. counting only those instructions that are actually executed at least once). We believe this approach provides a better insight on the actual instruction memory footprint and shows the difference between ISAs better. We found that these two number (our approach vs. measuring the size of the code) could be quite different for some applications since compilers might include source codes for all routines in an included library, while some of these routines may not be used at all.

To find static icount, we modified QEMU’s source code to add a new data structure to track and count unique PCs during the execution. These changes are made in a routine called EXEC\_CPU() which is used in all ISAs. To check the correctness of our model, for each ISA, we used several synthetic benchmarks and checked the QEMU’s output to the actual static icount computed manually.

**Dynamic Instruction Count:** Similar to static icount, dynamic icount is also an important metric to show the runtime behavior of ISAs. Dynamic icount can directly affect total runtime especially in simple in-order cores where Instruction Per Cycle (IPC) for these cores are mostly close to 1 thus the total runtime is determined by dynamic icount metric.

Dynamic icount is also computed in QEMU by modifying the source code to be able to count this metric during execution. The results are validated using a synthetic benchmark where number of iterations of a simple loop changed in different runs. We checked whether QEMU correctly reports dynamic icount as the number of iterations changed for each run.

**Per-PC Iteration Count:** Another interesting metric which used in this paper is per-PC iteration count where for each PC we report how many times this instruction has been executed. This could be very useful to find the hot regions in the code, and find the reason(s) behind why some applications have significantly higher/lower dynamic icount. More details will be shown in the Section ?? . QEMU is modified to count this metric too, during execution.

#### 3.2.2 gem5

gem5 is a well-known, open-source, cycle-accurate simulator which can simulate in-order and out-of-order cores, memory systems, and interconnect in details. Using gem5 enables us to find runtime statistics (e.g total

number of cycles) and micro-architecture related statistics (e.g. cache miss rate). gem5 supports ARM and very recently RISC-V. Unfortunately gem5 only supports an old version of MIPS and hence we removed the analysis for MIPS in gem5. To have a fair comparison, processor and memory system configurations (i.e. clock rate, issue width, cache levels and size, delays, etc.) are matched for both ARM and RISC-V processors. We use a simple single-issue, 4-stage pipeline with no prefetcher and a single-level cache as an in-order core for RISC-V and ARM, and use a more sophisticated 4-issue, out-of-order, with a direct/indirect prefetcher and two level caches as an out-of-order core in our experiments. Detailed for these two cores are shown in Table 2.

**Table 2: Simulation configuration for gem5.**

Parameters	In-order core	Out-of-Order core
Issue Width	1	4
Private Caches	I\$/D\$ 32KB	I\$/D\$ 32KB
Shared Caches	N/A	L2 256KB
Branch Predictor	N/A	Tournament BP and BTB
Prefetcher	N/A	stride/next-line prefetcher

**Cycle Count and IPC:** we use the number reported by gem5 for these metrics. Same input is used for both RISC-V and ARM and the numbers reported are from beginning to end of each application. We use IPC to show the effect of using different ISAs on processor’s computation speed. Our key findings about these two numbers will be shown in Section ??.

**Microarchitecture Statistics:** to gain some insights about the runtime behavior of each core and ISA’s impacts on them, we check several microarchitecture-related metrics such as: cache miss rate (MPKI), branch predictor accuracy, instruction mix, fetch and commit rates, total stall/idle/squashed cycles, memory bandwidth utilization, direct/indirect branches and function calls, and dependent memory load/stores across different applications.

### 3.2.3 McPAT

McPAT is an integrated power and area modeling simulator. It uses ITRS roadmap models at circuit level to model both static and dynamic power of the system. McPAT uses an XML-based interface to read microarchitecture-related statistics generated by a cycle-accurate simulator (e.g. gem5) as inputs and uses a detailed model for cores, memory system, NoC, etc. to estimate the area and power of the system. We use a set of python and shell scripts to parse the statistics from gem5 and fill in the XML template in McPAT.

**Dynamic Power and Total Energy:** using McPAT, we measure the dynamic power consumption of core (both in-order and out-of-order) and memory system. Using the data from gem5, we also calculate the overall energy for these ISAs.

## 3.3 Applications

We use a representative set of applications from a

standard open-source embedded system’s benchmark suite called MiBench. The MiBench suite is commonly used to evaluate the performance of processors intended for the embedded/IoT market, and it was designed to be representative of the computation that is needed in that market (e.g. automotive, industrial systems, etc.). These applications are mostly compute intensive and designed such that have minimum interactions with outside the processors. Many of the applications are also overlapped with EEMBC benchmark suite.

Our applications we picked ranges from basic math abilities (*basicmath*), bit manipulation (*bitcount*), simple data organization (*qsort*), a shape recognition program (*susan*), shortest path calculations (*dijkstra*), to data encryption, decryption and hashing (*blowfish*, *rijndael*, *sha*), and communications applications (*fft*, *crc32*, *adpcm*). In addition to these 11 applications picked from MiBench, we use two well-known open-source application which often used in industry to report the performance of the processor *Coremark* and *Dhrystone*. For each of these applications the “large” dataset is used. For Coremark and Dhrystone, iterations are chosen such that the total number of executed instructions be more than 200 million instructions (we chose 1000 iterations for Coremark and 1 million iterations for Dhrystone). For each application, same inputs are used for all runs across different ISAs/processors.

## 4. EXPERIMENTAL ANALYSIS

In this section we report our results and findings of static and dynamic instruction count, performance, power, and energy.

### 4.1 Instruction Count

Figure 1 shows the static icount obtained from QEMU. As seen in the figure, on average 32bit and 64bit ARM are about 15% more dense than that of in MIPS and RISC-V. Mibench applications have on average 5k instruction that are executed at least once. Figure 2 shows the dynamic icount obtained from QEMU. Unlike static icount, dynamic icount can be quite different from one application to another for different ISAs. However, on average MIPS, ARM, and RISC-V have almost same dynamic icount.

**Key Findings:** 1- Mixed/Combined instructions (e.g. add+shift, mult+add, etc.) and three operand/three-way comparison in ARM could result in significant dynamic and static icount reduction. A possible and interesting extension to RISC-V could be adding this sort of instructions to the base ISA (RV-G) for high-performance scenarios. An example is shown in Figure 3, where a same function is shown for ARM and RISC-V ISAs. As seen in this example, “ldr” instruction in ARM with embedded “lsl” instruction inside it, has saved one instruction. Further “cmn” (compare and add), also saved two extra instructions in ARM. We found that there are many examples such as this where more complex instructions in ARM could save more space, however, we will later show that this complexity comes with more power consumption.

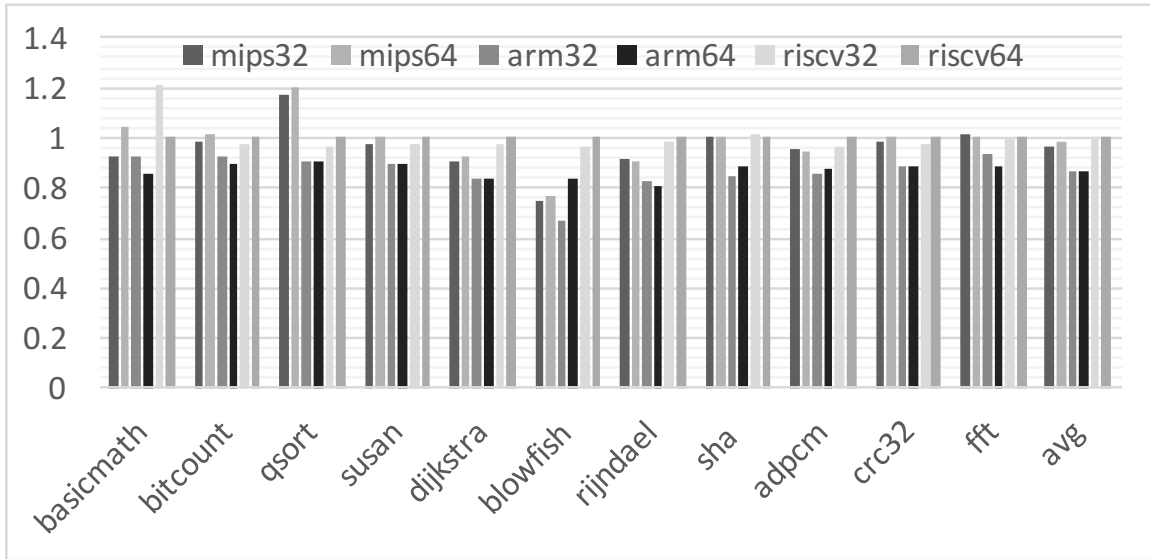


Figure 1: Static Instruction Count for MiBench applications using QEMU. Results are normalized w.r.t. RISC-V64. The average number of static instructions is about 5000 for these benchmarks.

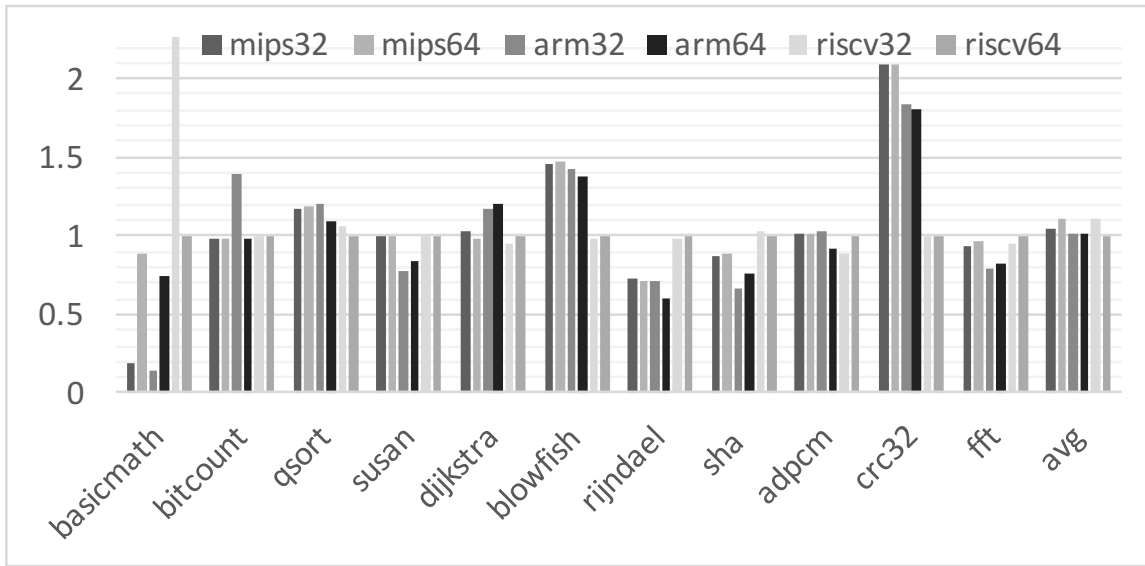
**Outliers:** static icount for qsort on MIPS is significantly higher. The main reason for that is due to the way a hot loop in a function called `MSORT_WITH_TEMP` (part of glibc library for quicksort) is implemented in MIPS where a few extra instructions are used for MIPS. Interestingly, we find that this function is implemented slightly different among different toolchains.

## 5. RELATED WORK

## 6. CONCLUSIONS

To best of our knowledge, this is first work to compare RISC-V with its popular proprietary counterparts (ARM and MIPS). We used state-of-art simulation and emulation frameworks: *qemu* for program analysis and *gem5* for microarchitectural simulations. We also present concrete cases where RISC-V clearly falls behind compared to ARM, MIPS ISAs and what addendum could possibly make it competitive. To our surprise, we also stumbled upon cases where RISC-V is better than proprietary ISAs. Overarching goal of our exploration is to enable RISC-V designers so that they can augment their designs and be able to close the gap with other state-of-art ISAs.

## 7. REFERENCES



**Figure 2: Dynamic Instruction Count for MiBench applications using QEMU. Results are normalized w.r.t. RISCv64. The average number of dynamic instructions for these benchmarks is about 450 million.**

### ARM:

```
400464: f8627aa0 ldr    x0, [x21, x2, lsl #3]
400468: 91000673 add     x19, x19, #0x1
40046c: ca5c201c eor     x28, x0, x28, lsr #8
400470: aa1403e0 mov     x0, x20
400474: 94002261 bl      408df8 <_IO_getc>
400478: ca000382 eor     x2, x28, x0
40047c: 3100041f cmn     w0, #0x1
400480: 92401c42 and     x2, x2, #0xff
400484: 54ffff01 b.ne   b.any 400464 <main+0x74> //
```

### RISCV:

```
10226: 6398 ld      a4,0(a5)
10228: 405 addi    s0,s0,1
1022a: 00d74d33 xor     s10,a4,a3
1022e: 8526 mv      a0,s1
10230: 039060ef jal     ra,16a68 <_IO_getc>
10234: 00ad47b3 xor     a5,s10,a0
10238: 0ff7f793 andi   a5,a5,255
1023c: 078e slli    a5,a5,0x3
1023e: 97ca add     a5,a5,s2
10240: 008d5693 srli    a3,s10,0x8
10244: ff3511e3 bne     a0,s3,10226<main+0x66>
```

Figure 3: A code snippet in assembly showing a same function from Mibench benchmark suite for ARM-64 and RISCV-64 ISAs. Differences shown in rectangles.